

AiLibi — Task 18.26 Owner Brief

The real-LLM finalist eval · PR #317 merged 2026-08-01 · baseline-6 substrate · Qwen/Qwen3.6-27B via Featherless · seeds 0-49 · \$0

1 · Campaign at a glance

Nine arms recorded on the real path at 50 seeds each (449 kept games; one documented exclusion): four impostor finalists, a fresh all-scripted FSM comparator, and four crew diagnostics against the frozen champion. Pre-registration (both cells, the slate, the protocol) was committed before any seed ran. Full attempt wall 105.2 h over 505 attempts at the rolling two-leg posture (12.5 min per recorded game — 2.5% over the pre-registered rate). Every measurement, split view, derived comparison cell, and duration digest is persisted in **training/reports/results-finalist-eval.jsonl** and guarded by 31 pin tests; the 17.14 rows are preserved untouched. The evidence survived a 17-round automated review (~45 comments: fixed, rebutted from the ratified texts, or owner-ruled).

2 · Axis 1 — the selection table

arm	role	win rate	delta vs FSM	referee	flags (floor 1.091)	conversion (derived floor)	witnessed (floor 0.034)
ea4bc955	champion (slate 1)	0.52	+0.26	FAIL	0.936	0.367 / 0.669	0.152
bfd145cb	runner-up (slate 2)	0.56	+0.30	FAIL	0.900	0.351 / 0.695	0.148
6d327dcb	incumbent (slate 3)	0.38	+0.12	FAIL	0.969	0.444 / 0.646	0.223
7f73929d	F13 arm (slate 4) - n=49	0.429	+0.184 (intersection)	FAIL	0.828	0.389 / 0.755	0.220
FSM comparator	all-scripted	0.26	-	PASS	1.198	0.552 / 0.523	0.046

The one-line story: every learned arm beats the scripted baseline on wins, and every learned arm fails the referee — the comparator is the slate's only referee PASS. The evidence-supply-for-wins trade is real at slate level, and win rank inverts gauge rank (bfd145cb wins most with the worst gauges; the incumbent wins least with the best). Note the substrate cost the incumbent 14 points (0.52 at baseline-5 to 0.38 here). All validity gates PASS on the five arms above; witnessed_event_rate is UNRESOLVABLE everywhere by the pre-registered noise rule and feeds no ruling.

3 · The F13 cell (champions vs runner-ups) — answered, awaiting your ruling

Pre-registered hypotheses: **A** — the ES trades evidence-supply for wins and runner-ups sit one step less far along the trade (their screening-time gauge margins persist at n=50); **B** — the n<=6 referee reads were noise (the gap vanishes). Measured on the common 49-seed intersection:

gauge	champion mean	runner-up mean	margin	pooled-side noises	barred from supporting A by
witnessed_event_rate	0.192	0.185	-0.0065	0.070 / 0.060	both sides' noise
flags_per_meeting	0.955	0.863	-0.0919	0.180 / 0.157	both sides' noise
testimony_backed_conv.	0.404	0.369	-0.0348	0.066 / 0.009	champion side's noise

No gauge supports hypothesis A under the registered either-side noise rule, and all three margins point the wrong way for A (negative — champions closer to the floors, not runner-ups). The shape is hypothesis-B (screening reads were noise), with the guard that A-unsupported is not B-demonstrated. The within-lineage pair (ea4bc955 vs bfd145cb) collapses to one readable gauge (conversion, -0.022). **The ruling is yours in 18.27.**

4 · Axis 2 — the emergence picture (the campaign's strongest finding)

The clause-(a) statistics (pooled two-proportion z vs the same-seed FSM comparator, |z| at or over 1.96 starred) split the instrument set cleanly along the physical/language line:

registered cell	ea4bc955	bfd145cb	6d327dcb	7f73929d	arms over bar
crew-witnessed kill rate	+3.37 *	+3.27 *	+4.89 *	+4.77 *	4 / 4
co-present departure	+4.32 *	+4.37 *	+6.00 *	+5.73 *	4 / 4
false-vouch saw_player	+0.76	+1.13	+2.45 *	+1.32	1
false-vouch corroboration	+1.55	+1.11	+3.81 *	+2.77 *	2
frame attempts / conv. / alibi / deflection	under bar	under bar	under bar	under bar	0
one-hop point-biserial (Fisher)	-0.65	-0.68	+2.81 *	+0.75	1

Reading: the ES learned *physical kill-craft*, not rhetoric. Both physical cells clear the bar on all four arms — and co-present killing is literally absent from scripted play (corpus 0/505, comparator 0/174, every learned arm 0.076-0.187). The language layer is flat except the incumbent 6d327dcb, the lone deception outlier (5 cells over bar). Each star is one limb of four: the split-reproducibility views are committed per row (clause b), zero of the five campaign ablation runs are clause-(c)-complete by their own reports, and both entropy rulings are NOT-DEMONSTRATED (the variance field's routed contract never landed). **No full emergence claim can fire this phase** — what you have is strong, honest observation evidence.

5 · Crew diagnostics and the kill-rate rider

c1 (owned-tasks lineage) plays; c2 (general base) starves. c1-gen9 is the only clean crew arm (validity PASS, meeting rate 1.0, 26/50 vs the champion) but the paired same-seed read against its gen-0 control is a null: 38/49 concordant seeds, discordants split 6-5, McNemar p = 1.0 — no learning signal in either direction. c2-gen0 is total starvation (zero meetings, zero LLM calls, crew 1/50); c2-gen9 reached exactly the 0.60 floor and produced two LLM-free deterministic stalemates — meeting-triggering is the one thing that lineage partially learned. **The routed 18.25 rider is answered: the elevated crew-witnessed kill rate is an artifact of learned-impostor kill placement** (gen-9 and gen-0 within a point of each other at the same opponent; elevation present against scripted crew and in the LLM-free leg; formally inconclusive as an equivalence — no margin was pre-registered). All crew axis-2 cells read NOT-DEMONSTRABLE (your ruling: no opponent-matched scripted-crew comparator was recorded).

6 · Anomalies and gotchas worth remembering

- **Seed 35 (7f73929d) is unrecordable at this substrate:** 14 attempts, every one the identical validation default at the first meeting's opening turn (deterministic pre-meeting prefix = identical prompt; the model reliably emits a schema-invalid turn). Slot 4 ships at n=49 with intersection cells for every affected read. The corresponding DoD box is explicitly un-ticked in the PR.
- **witnessed_event_rate cannot be resolved at n=50** under the 25% noise precondition — on any arm. If a future protocol needs that gauge to rule, it needs more seeds, not a waiver.
- **The corpus JSON has a trap:** the subkey named baseline_cells_corpus_9p2i is a stale baseline-5 snapshot (byte-identical across both roster entries); the true baseline-6 cells are the top-level blocks. Worth a cleanup task so 18.27+ tooling never quotes the wrong block.
- **Three stuck-seed classes now documented:** retryable validation-defaults (sampling variance), LLM-free deterministic stalemates (unretryable by construction), and meeting-bearing robust stalemates. Classes 2 and 3 exit rc 0 — record-time retry machinery never sees them.

- **Operational:** a failed attempt costs the same wall as a successful one (12.5 min); the two starved c2 legs ran at engine speed (50 games in 27 seconds of summed wall for c2-gen0).

7 · What you must rule for continuation (the 18.27 owner checklist)

18.27 is an OWNER task — these rulings are yours, then the memo + the ruled branch get implemented (this thread has full context, or dispatch the generated prompt with your rulings attached).

#	ruling	the evidence in front of you	branches
1	The flip (axis 1): does the champion clear the §1.3 bar?	ea4bc955: win 0.52, +0.26 over the same-seed FSM row, validity PASS — but referee FAIL (flags -0.155, conversion -0.302) with the witnessed gauge excluded as UNRESOLVABLE. Read the bar's verbatim terms in audit-phase-17-close.md §1.3 against §16.a.	PASS: artifact surface swap + pre-author the 18.28 selector flip (default provably not yet moved). FAIL: default provably unmoved.
2	F13: hypothesis A or B?	No gauge supports A (either-side noise rule + direction); B-shaped but not demonstrated. Also decide what this means for trusting small-n referee screens in future campaigns.	Rule A / rule B / rule insufficient.
3	Emergence claims (axis 2): what status do the physical-layer observations get?	Witnessed + co-present significant on 4/4 arms vs same-seed FSM; splits committed; but zero clause-(c)-complete ablations and entropy NOT-DEMONSTRATED — the §6 conjunction cannot complete.	Observations-only this phase, or route an ablation follow-up contract to complete clause (c).
4	Crew adoption: confirm none.	18.25 named no finalist; 18.26 concurs (c1 null vs control, c2 starves).	Confirm; optionally route the scripted-crew-vs-champion comparator arm (~10 h) if crew claims are ever wanted.
5	Follow-up contracts to route (or drop), explicitly:	(a) entropy-variance field plumbing (pre-req for any future entropy ruling); (b) the scripted-crew-vs-ea4bc955 comparator arm; (c) the corpus-JSON legacy-block cleanup; (d) seed-35 re-record under a future substrate if slot 4 must reach n=50.	Route / drop each, in the 18.27 memo.
6	Housekeeping:	~/ailibi-campaign-1826/ holds the recordings, leg logs, and the only copies of the seed-35 forensics — keep until 18.27 lands, then reclaimable.	Keep / archive.

8 · Open questions (non-blocking, flagged for your read)

- **Champion identity:** bfd145cb out-wins the ratified champion (0.56 vs 0.52) with worse gauges. The slate ratification named ea4bc955 the champion arm; if 18.27's bar reading turns on win edge alone, does the runner-up's edge matter, or is that exactly the trade F13 warns about?
- **Off-menu is vacuous everywhere** (0 in ~6,600 decisions per arm, corpus included). Perfect compliance, or an instrument that can no longer trigger on this substrate? Cheap to check once, worth knowing before it is quoted as a safety result.
- **The 6d327dcb deception outlier:** the incumbent — trained longest under the old substrate — is the only arm with language-layer significance. Artifact of its lineage, or the earliest visible rhetoric learning? The ablation route in ruling 3 would answer it.

Sources: training/reports/report-finalist-eval.md Part II (§8-§17) and results-finalist-eval.jsonl at merge commit of PR #317; pins in tests/training/test_finalist_eval_pins.py. This brief is a personal summary, not a repo artifact — every number above is quoted from the committed evidence.